



Scaling Behavior of H-Net Tokenization

(a) Compression vs Training Compute



(b) Tokenization Efficiency vs Compute



(c) Improvements Relative to 1-Epoch Baseline

